

▼ Libraries

```
import os
from os import listdir

import pandas as pd
import numpy as np
from numpy import hstack
import matplotlib.pyplot as plt
from matplotlib.image import imread
import plotly.express as px
import seaborn as sns

import tensorflow as tf
from keras.utils import to_categorical
from keras.preprocessing import image
from keras.layers import Dense, Flatten, Activation, Dropout
from keras import optimizers
```

▼ Raw Data Import

```
from google.colab import drive
drive.mount('/content/drive/')

!ls /content/drive/

Mounted at /content/drive/
MyDrive
```

▼ Data Preprocessing

```
dicom_info = pd.read_csv('/content/drive/MyDrive/CIND860/csv/dicom_info.csv')
mass_case_test = pd.read_csv('/content/drive/MyDrive/CIND860/csv/mass_case_description_test_set.csv')
mass_case_train = pd.read_csv('/content/drive/MyDrive/CIND860/csv/mass_case_description_train_set.csv')
image_dir = '/content/drive/MyDrive/CIND860/jpeg'

full_mammogram_images = dicom_info[dicom_info.SeriesDescription == 'full mammogram images'].image_path
cropped_images = dicom_info[dicom_info.SeriesDescription == 'cropped images'].image_path
roi_mask_images = dicom_info[dicom_info.SeriesDescription == 'ROI mask images'].image_path

full_mammogram_images = full_mammogram_images.apply(lambda x: x.replace('CBIS-DDSM/jpeg', image_dir))
cropped_images = cropped_images.apply(lambda x: x.replace('CBIS-DDSM/jpeg', image_dir))
roi_mask_images = roi_mask_images.apply(lambda x: x.replace('CBIS-DDSM/jpeg', image_dir))
full_mammogram_images.iloc[0]

'/content/drive/MyDrive/CIND860/jpeg/1.3.6.1.4.1.9590.100.1.2.2483867420106785823
090005272213277814849/1-249.jpg'
```

Preparing the image paths for preprocessing

```
full_mammogram_dict = dict()
cropped_dict = dict()
roi_mask_dict = dict()

for dicom in full_mammogram_images:
    key = dicom.split("/")[6]
    full_mammogram_dict[key] = dicom
for dicom in cropped_images:
    key = dicom.split("/")[6]
    cropped_dict[key] = dicom
for dicom in roi_mask_images:
    key = dicom.split("/")[6]
    roi_mask_dict[key] = dicom
next(iter((full_mammogram_dict.items())))
```

```
('1.3.6.1.4.1.9590.100.1.2.248386742010678582309005372213277814849',
 '/content/drive/MyDrive/CIND860/jpeg/1.3.6.1.4.1.9590.100.1.2.248386742010678582309005372213277814849/1-249.jpg')
```

```
def fix_image_path(dataset):
    for i, img in enumerate(dataset.values):
        img_name = img[11].split("/")[2]
        dataset.iloc[i,11] = full_mammogram_dict[img_name]
        img_name = img[12].split("/")[2]
        dataset.iloc[i,12] = cropped_dict[img_name]
        img_name = img[13].split("/")[2]
        dataset.iloc[i,13] = roi_mask_dict[img_name]

fix_image_path(mass_case_train)
fix_image_path(mass_case_test)
mass_data = mass_case_train.append(mass_case_test)
mass_data.head()

<ipython-input-7-5d9cb52cd42a>:3: FutureWarning: The frame.append method is deprecated
mass_data = mass_case_train.append(mass_case_test)
```

patient_id	breast_density	left		image	abnormality	abnormality			
		right	or				view	id	type
			breast						
0	P_00001	3	LEFT	CC	1	mass ARCHITECT			
1	P_00001	3	LEFT	MLO	1	mass ARCHITECT			
2	P_00004	3	LEFT	CC	1	mass ARCHITECT			
3	P_00004	3	LEFT	MLO	1	mass ARCHITECT			
4	P_00004	3	RIGHT	MLO	1	mass			

▼ Data Cleaning

```
dicom_info.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10237 entries, 0 to 10236
Data columns (total 38 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   file_path        10237 non-null   object 
 1   image_path       10237 non-null   object 
 2   AccessionNumber  10237 non-null   float64
 3   BitsAllocated   10237 non-null   int64  
 4   BitsStored      10237 non-null   int64  
 5   BodyPartExamined 10237 non-null   object 
 6   Columns          10237 non-null   int64  
 7   ContentDate     10237 non-null   int64  
 8   ContentTime      10237 non-null   float64
 9   ConversionType  10237 non-null   object 
 10  HighBit         10237 non-null   int64  
 11  InstanceNumber  10237 non-null   int64  
 12  LargestImagePixelValue 10237 non-null   int64  
 13  Laterality      9671 non-null    object 
 14  Modality        10237 non-null   object 
 15  PatientBirthDate 0 non-null     float64
 16  PatientID       10237 non-null   object 
 17  PatientName     10237 non-null   object 
 18  PatientOrientation 10237 non-null   object 
 19  PatientSex      0 non-null     float64
 20  PhotometricInterpretation 10237 non-null   object 
 21  PixelRepresentation 10237 non-null   int64  
 22  ReferringPhysicianName 0 non-null     float64
 23  Rows             10237 non-null   int64  
 24  SOPClassUID     10237 non-null   object 
 25  SOPInstanceUID  10237 non-null   object 
 26  SamplesPerPixel 10237 non-null   int64  
 27  SecondaryCaptureDeviceManufacturer 10237 non-null   object 
 28  SecondaryCaptureDeviceManufacturerModelName 10237 non-null   object 
 29  SeriesDescription 9671 non-null    object 
 30  SeriesInstanceUID 10237 non-null   object
```

```

31 SeriesNumber          10237 non-null  int64
32 SmallestImagePixelValue 10237 non-null  int64
33 SpecificCharacterSet   10237 non-null  object
34 StudyDate             9671 non-null   float64
35 StudyID               10237 non-null  object
36 StudyInstanceUID      10237 non-null  object
37 StudyTime              9671 non-null   float64
dtypes: float64(7), int64(12), object(19)
memory usage: 3.0+ MB

```

```

dicom_info.drop(['PatientBirthDate', 'AccessionNumber', 'Columns', 'ContentDate', 'ContentTime', 'PatientSex', 'PatientBirthDate',
                 'ReferringPhysicianName', 'Rows', 'SOPClassUID', 'SOPInstanceUID',
                 'StudyDate', 'StudyID', 'StudyInstanceUID', 'StudyTime', 'InstanceNumber', 'SeriesInstanceUID', 'S

```

```
dicom_info.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10237 entries, 0 to 10236
Data columns (total 21 columns):
 #   Column           Non-Null Count Dtype  
 --- 
 0   file_path        10237 non-null  object  
 1   image_path       10237 non-null  object  
 2   BitsAllocated    10237 non-null  int64   
 3   BitsStored       10237 non-null  int64   
 4   BodyPartExamined 10237 non-null  object  
 5   ConversionType   10237 non-null  object  
 6   HighBit          10237 non-null  int64   
 7   LargestImagePixelValue 10237 non-null  int64  
 8   Laterality        9671 non-null   object  
 9   Modality          10237 non-null  object  
 10  PatientID         10237 non-null  object  
 11  PatientName       10237 non-null  object  
 12  PatientOrientation 10237 non-null  object  
 13  PhotometricInterpretation 10237 non-null  object  
 14  PixelRepresentation 10237 non-null  int64   
 15  SamplesPerPixel    10237 non-null  int64   
 16  SecondaryCaptureDeviceManufacturer 10237 non-null  object  
 17  SecondaryCaptureDeviceManufacturerModelName 10237 non-null  object  
 18  SeriesDescription   9671 non-null   object  
 19  SmallestImagePixelValue 10237 non-null  int64  
 20  SpecificCharacterSet 10237 non-null  object  
dtypes: int64(7), object(14)
memory usage: 1.6+ MB

```

```
dicom_info.isna().sum()
```

file_path	0
image_path	0
BitsAllocated	0
BitsStored	0
BodyPartExamined	0
ConversionType	0
HighBit	0
LargestImagePixelValue	0
Laterality	566
Modality	0
PatientID	0
PatientName	0
PatientOrientation	0
PhotometricInterpretation	0
PixelRepresentation	0
SamplesPerPixel	0
SecondaryCaptureDeviceManufacturer	0
SecondaryCaptureDeviceManufacturerModelName	0
SeriesDescription	566
SmallestImagePixelValue	0
SpecificCharacterSet	0

dtype: int64

```

dicom_info['SeriesDescription'].fillna(method = 'bfill', axis = 0, inplace=True)
dicom_info['Laterality'].fillna(method = 'bfill', axis = 0, inplace=True)

```

```
dicom_info.isna().sum()
```

file_path	0
image_path	0
BitsAllocated	0
BitsStored	0

```
BodyPartExamined          0
ConversionType            0
HighBit                   0
LargestImagePixelValue    0
Laterality                 0
Modality                  0
PatientID                 0
PatientName                0
PatientOrientation          0
PhotometricInterpretation    0
PixelRepresentation          0
SamplesPerPixel             0
SecondaryCaptureDeviceManufacturer 0
SecondaryCaptureDeviceManufacturerModelName 0
SeriesDescription           0
SmallestImagePixelValue     0
SpecificCharacterSet        0
dtype: int64
```

```
dicom_info['SeriesDescription'].value_counts()
```

```
cropped images      3859
ROI mask images    3340
full mammogram images 3038
Name: SeriesDescription, dtype: int64
```

```
mass_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 1696 entries, 0 to 377
Data columns (total 14 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   patient_id       1696 non-null    object  
 1   breast_density    1696 non-null    int64  
 2   left or right breast 1696 non-null    object  
 3   image view        1696 non-null    object  
 4   abnormality id    1696 non-null    int64  
 5   abnormality type   1696 non-null    object  
 6   mass shape         1692 non-null    object  
 7   mass margins        1636 non-null    object  
 8   assessment         1696 non-null    int64  
 9   pathology          1696 non-null    object  
 10  subtlety          1696 non-null    int64  
 11  image file path    1696 non-null    object  
 12  cropped image file path 1696 non-null    object  
 13  ROI mask file path 1696 non-null    object  
dtypes: int64(4), object(10)
memory usage: 198.8+ KB
```

```
mass_data['image view'] = mass_data['image view'].astype('category')
mass_data['mass margins'] = mass_data['mass margins'].astype('category')
mass_data['mass shape'] = mass_data['mass shape'].astype('category')
mass_data['abnormality type'] = mass_data['abnormality type'].astype('category')
mass_data['pathology'] = mass_data['pathology'].astype('category')
mass_data.isna().sum()
```

```
patient_id          0
breast_density       0
left or right breast 0
image view          0
abnormality id       0
abnormality type     0
mass shape           4
mass margins          60
assessment           0
pathology             0
subtlety              0
image file path      0
cropped image file path 0
ROI mask file path    0
dtype: int64
```

```
mass_data['mass shape'].fillna(method = 'bfill', axis = 0, inplace=True)
mass_data['mass margins'].fillna(method = 'bfill', axis = 0, inplace=True)
```

```
mass_data.isna().sum()
```

```

patient_id          0
breast_density      0
left or right breast 0
image view          0
abnormality id      0
abnormality type    0
mass shape           0
mass margins          0
assessment            0
pathology              0
subtlety                0
image file path       0
cropped image file path 0
ROI mask file path     0
dtype: int64

```

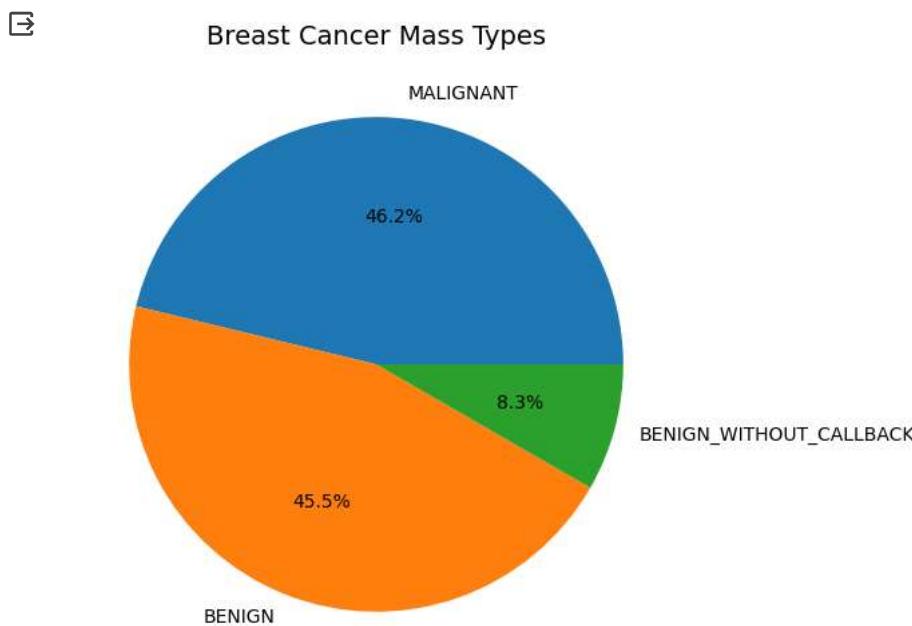
▼ Data Visualization

```

value = mass_data['pathology'].value_counts()
plt.figure(figsize=(8,6))

plt.pie(value, labels=value.index, autopct='%1.1f%%')
plt.title('Breast Cancer Mass Types', fontsize=14)
plt.show()

```



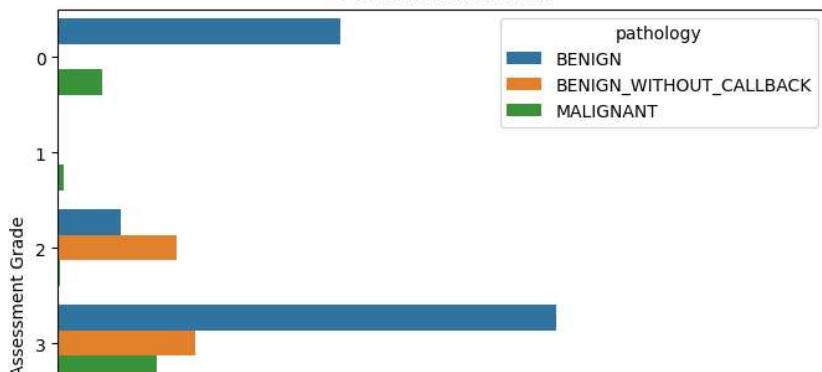
```

plt.figure(figsize=(8,6))
sns.countplot(mass_data, y='assessment', hue='pathology')
plt.title('Breast Cancer Assessment\n\n 0: Undetermined || 1: Well Differentiated\n 2: Moderately differentiated || 3: Poorly Differentiated')
plt.xlabel('Count')
plt.ylabel('Assessment Grade')
plt.show()

```

Breast Cancer Assessment

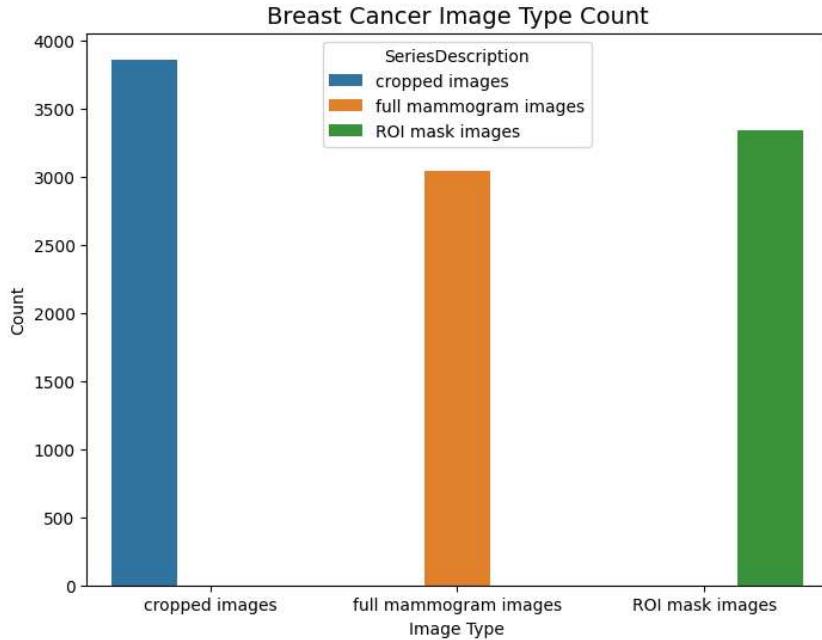
0: Undetermined || 1: Well Differentiated
 2: Moderately differentiated || 3: Poorly Differentiated
 4-5: Undifferentiated



```
images = dicom_info['SeriesDescription'].value_counts()
images = images.reset_index()
images = images.rename(columns={'index':'SeriesDescription','SeriesDescription':'counts'})

images = dicom_info['SeriesDescription'].value_counts()
plt.figure(figsize=(8,6))

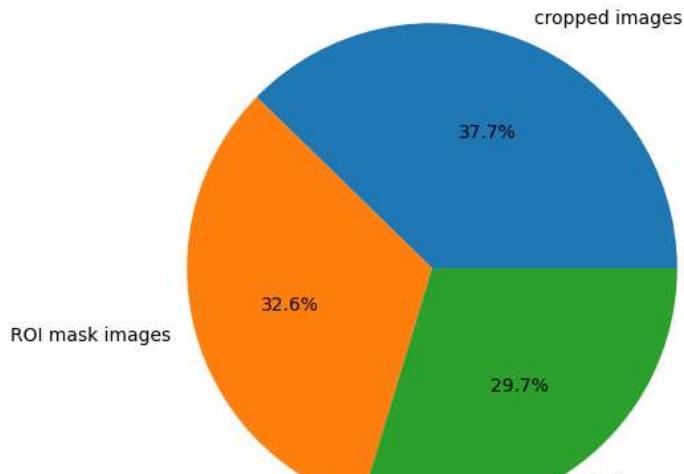
sns.countplot(dicom_info, x='SeriesDescription', hue='SeriesDescription', orient='h')
plt.title('Breast Cancer Image Type Count', fontsize=14)
plt.xlabel('Image Type')
plt.ylabel('Count')
plt.show()
```



```
images = dicom_info['SeriesDescription'].value_counts()
plt.figure(figsize=(8,6))

plt.pie(images, labels=images.index, autopct='%1.1f%%')
plt.title('Breast Cancer Image Types', fontsize=14)
plt.show()
```

Breast Cancer Image Types



▼ Image Enhancement

```

import tensorflow as tf
import cv2
from sklearn.model_selection import train_test_split
from tensorflow.keras.utils import to_categorical
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.utils import plot_model

def image_processor(image_path, target_size):
    absolute_image_path = os.path.abspath(image_path)
    image = cv2.imread(absolute_image_path)
    image = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)
    image = cv2.resize(image, (target_size[1], target_size[0]))
    image_array = image / 255.0
    return image_array

# Define the target size
target_size = (224, 224, 3)

# Apply preprocessor to train data
mass_data['processed images'] = mass_data['image file path'].apply(lambda x: image_processor(x, target_size))

# Create a binary mapper
class_mapper = {'MALIGNANT': 1, 'BENIGN': 0, 'BENIGN_WITHOUT_CALLBACK': 0}

# Convert the processed_images column to an array
X_resized = np.array(mass_data['processed images'].tolist())

# Apply class mapper to pathology column
mass_data['labels'] = mass_data['pathology'].replace(class_mapper)

# Check the number of classes
num_classes = len(mass_data['labels'].unique())

```

▼ Preparing the Train/Test Set

```

# Split data into train, test, and validation sets
X_train, X_temp, y_train, y_temp = train_test_split(X_resized, mass_data['labels'].values, test_size=0.3, random_state=42)
X_test, X_val, y_test, y_val = train_test_split(X_temp, y_temp, test_size=0.33, random_state=42)

# Convert integer labels to one-hot encoded labels
y_train = to_categorical(y_train, num_classes)

```

```

y_test = to_categorical(y_test, num_classes)
y_val = to_categorical(y_val, num_classes)

print("X_train: ", X_train.shape)
print("y_train: ", y_train.shape)
print("X_test: ", X_test.shape)
print("y_test: ", y_test.shape)
print("X_val: ", X_val.shape)
print("y_val: ", y_val.shape)

X_train: (1187, 224, 224, 3)
y_train: (1187, 2)
X_test: (341, 224, 224, 3)
y_test: (341, 2)
X_val: (168, 224, 224, 3)
y_val: (168, 2)

# Convert one-hot encoded labels to single values for CNN Model
y_train = np.argmax(y_train, axis=1)
y_val = np.argmax(y_val, axis=1)
y_test = np.argmax(y_test, axis=1)

print("y_train: ", y_train.shape)
print("y_test: ", y_test.shape)
print("y_val: ", y_val.shape)

y_train: (1187,)
y_test: (341,)
y_val: (168,)

```

▼ CNN Model

```

from keras.models import Sequential
from keras import layers, models
from keras.optimizers import RMSprop, Adam
from keras.layers import Flatten, Dense, Embedding

model0 = models.Sequential()

model0.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(224, 224, 3)))
model0.add(layers.MaxPooling2D((2, 2)))
model0.add(layers.Conv2D(64, (3, 3), activation='relu'))
model0.add(layers.MaxPooling2D((2, 2)))
model0.add(layers.Conv2D(128, (3, 3), activation='relu'))
model0.add(layers.MaxPooling2D((2, 2)))

model0.add(layers.Flatten())
model0.add(layers.Dense(48, activation='relu'))
model0.add(layers.Dropout(0.5))
model0.add(layers.Dense(1, activation='sigmoid'))

model0.summary()

```

Model: "sequential"

Layer (type)	Output Shape	Param #
<hr/>		
conv2d (Conv2D)	(None, 222, 222, 32)	896
<hr/>		
max_pooling2d (MaxPooling2D)	(None, 111, 111, 32)	0
conv2d_1 (Conv2D)	(None, 109, 109, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 54, 54, 64)	0
conv2d_2 (Conv2D)	(None, 52, 52, 128)	73856
max_pooling2d_2 (MaxPooling2D)	(None, 26, 26, 128)	0
flatten (Flatten)	(None, 86528)	0

```

dense (Dense)           (None, 48)          4153392
dropout (Dropout)        (None, 48)          0
dense_1 (Dense)         (None, 1)           49
=====
Total params: 4246689 (16.20 MB)
Trainable params: 4246689 (16.20 MB)
Non-trainable params: 0 (0.00 Byte)

```

```

model0.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])

history0 = model0.fit(X_train, y_train, epochs=50, batch_size=32, validation_data=(X_val, y_val))

test_loss, test_acc = model0.evaluate(X_test, y_test)
print(f'Test accuracy: {test_acc}')

Epoch 23/50
38/38 [=====] - 2s 51ms/step - loss: 0.1186 - accuracy: 0.9419 - val_loss: 2.0542 - val_accuracy: 0.6071
Epoch 24/50
38/38 [=====] - 2s 51ms/step - loss: 0.0883 - accuracy: 0.9612 - val_loss: 2.6212 - val_accuracy: 0.6071
Epoch 25/50
38/38 [=====] - 2s 52ms/step - loss: 0.0891 - accuracy: 0.9562 - val_loss: 2.2166 - val_accuracy: 0.6250
Epoch 26/50
38/38 [=====] - 2s 55ms/step - loss: 0.1027 - accuracy: 0.9570 - val_loss: 2.2684 - val_accuracy: 0.6310
Epoch 27/50
38/38 [=====] - 2s 56ms/step - loss: 0.1117 - accuracy: 0.9461 - val_loss: 1.8412 - val_accuracy: 0.6548
Epoch 28/50
38/38 [=====] - 2s 54ms/step - loss: 0.0829 - accuracy: 0.9638 - val_loss: 2.2334 - val_accuracy: 0.6369
Epoch 29/50
38/38 [=====] - 2s 51ms/step - loss: 0.1079 - accuracy: 0.9537 - val_loss: 2.1398 - val_accuracy: 0.5952
Epoch 30/50
38/38 [=====] - 2s 50ms/step - loss: 0.1133 - accuracy: 0.9419 - val_loss: 2.0230 - val_accuracy: 0.6369
Epoch 31/50
38/38 [=====] - 2s 51ms/step - loss: 0.0860 - accuracy: 0.9604 - val_loss: 2.4073 - val_accuracy: 0.6190
Epoch 32/50
38/38 [=====] - 2s 52ms/step - loss: 0.0923 - accuracy: 0.9511 - val_loss: 2.2938 - val_accuracy: 0.6548
Epoch 33/50
38/38 [=====] - 2s 54ms/step - loss: 0.1015 - accuracy: 0.9520 - val_loss: 2.0950 - val_accuracy: 0.6667
Epoch 34/50
38/38 [=====] - 2s 56ms/step - loss: 0.0950 - accuracy: 0.9537 - val_loss: 2.1845 - val_accuracy: 0.6548
Epoch 35/50
38/38 [=====] - 2s 57ms/step - loss: 0.0822 - accuracy: 0.9587 - val_loss: 2.4846 - val_accuracy: 0.6607
Epoch 36/50
38/38 [=====] - 2s 54ms/step - loss: 0.0740 - accuracy: 0.9621 - val_loss: 2.4416 - val_accuracy: 0.6429
Epoch 37/50
38/38 [=====] - 2s 51ms/step - loss: 0.0753 - accuracy: 0.9638 - val_loss: 2.5250 - val_accuracy: 0.6548
Epoch 38/50
38/38 [=====] - 2s 50ms/step - loss: 0.0639 - accuracy: 0.9722 - val_loss: 2.8634 - val_accuracy: 0.6310
Epoch 39/50
38/38 [=====] - 2s 52ms/step - loss: 0.0635 - accuracy: 0.9638 - val_loss: 2.3286 - val_accuracy: 0.6310
Epoch 40/50
38/38 [=====] - 2s 51ms/step - loss: 0.0635 - accuracy: 0.9638 - val_loss: 2.8524 - val_accuracy: 0.6488
Epoch 41/50
38/38 [=====] - 2s 51ms/step - loss: 0.0548 - accuracy: 0.9714 - val_loss: 2.7208 - val_accuracy: 0.5952
Epoch 42/50
38/38 [=====] - 2s 56ms/step - loss: 0.0575 - accuracy: 0.9646 - val_loss: 2.9406 - val_accuracy: 0.6607
Epoch 43/50
38/38 [=====] - 2s 59ms/step - loss: 0.0599 - accuracy: 0.9655 - val_loss: 2.9115 - val_accuracy: 0.6190
Epoch 44/50
38/38 [=====] - 2s 54ms/step - loss: 0.0690 - accuracy: 0.9680 - val_loss: 2.4528 - val_accuracy: 0.6429
Epoch 45/50
38/38 [=====] - 2s 51ms/step - loss: 0.0696 - accuracy: 0.9621 - val_loss: 3.0315 - val_accuracy: 0.6310
Epoch 46/50
38/38 [=====] - 2s 52ms/step - loss: 0.0693 - accuracy: 0.9604 - val_loss: 2.6472 - val_accuracy: 0.6488
Epoch 47/50
38/38 [=====] - 2s 51ms/step - loss: 0.0735 - accuracy: 0.9537 - val_loss: 2.9675 - val_accuracy: 0.6310
Epoch 48/50
38/38 [=====] - 2s 51ms/step - loss: 0.0613 - accuracy: 0.9714 - val_loss: 2.8287 - val_accuracy: 0.6548
Epoch 49/50
38/38 [=====] - 2s 52ms/step - loss: 0.0693 - accuracy: 0.9604 - val_loss: 2.6472 - val_accuracy: 0.6488
Epoch 50/50
38/38 [=====] - 2s 58ms/step - loss: 0.0711 - accuracy: 0.9655 - val_loss: 3.3164 - val_accuracy: 0.5833
11/11 [=====] - 1s 65ms/step - loss: 3.3923 - accuracy: 0.6334
Test accuracy: 0.633431077003479

```

```

acc = history0.history['accuracy']
val_acc = history0.history['val_accuracy']
loss = history0.history['loss']

```

```

val_loss = history0.history['val_loss']

epochs = range(1, len(acc) + 1)

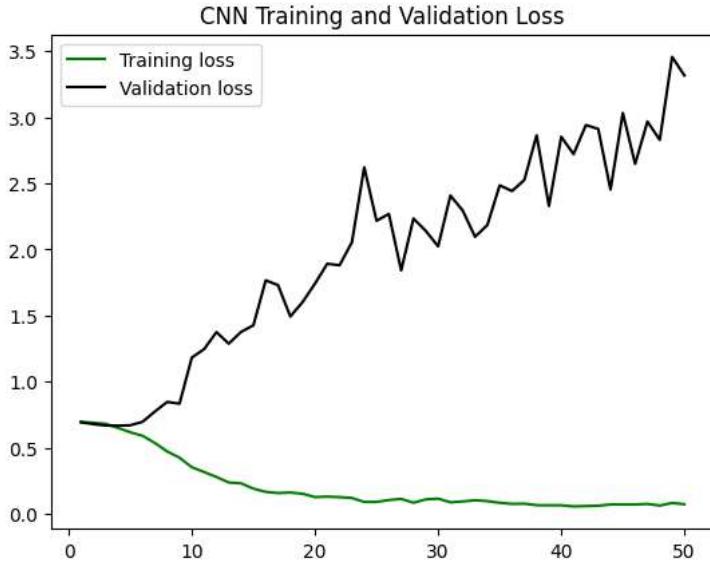
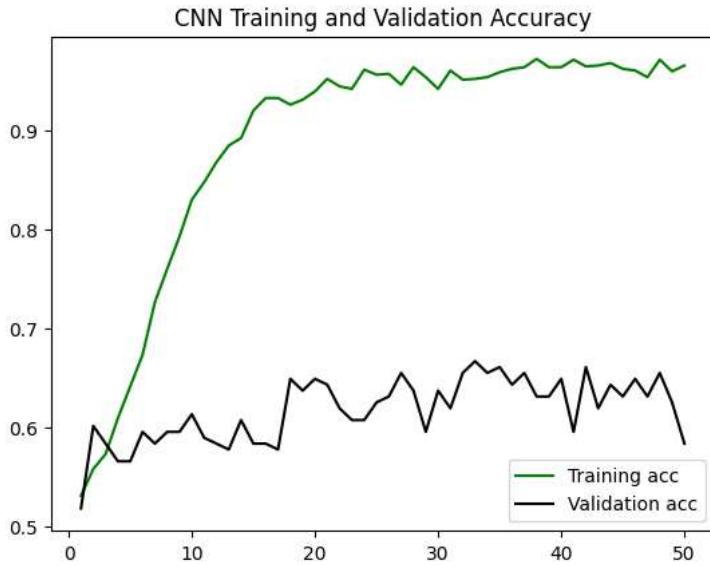
plt.plot(epochs, acc, 'g', label='Training acc')
plt.plot(epochs, val_acc, 'k', label='Validation acc')
plt.title('CNN Training and Validation Accuracy')
plt.legend()

plt.figure()

plt.plot(epochs, loss, 'g', label='Training loss')
plt.plot(epochs, val_loss, 'k', label='Validation loss')
plt.title('CNN Training and Validation Loss')
plt.legend()

plt.show()

```



```
from sklearn.metrics import classification_report, confusion_matrix
```

```

y_pred = model0.predict(X_test)
y_pred_classes = np.argmax(y_pred, axis=1)
y_true_classes = y_test

report0 = classification_report(y_true_classes, y_pred_classes)
print("Classification Report:")
print(report0)

cm0 = confusion_matrix(y_true_classes, y_pred_classes)

```

```

print("Confusion Matrix:")
print(cm0)

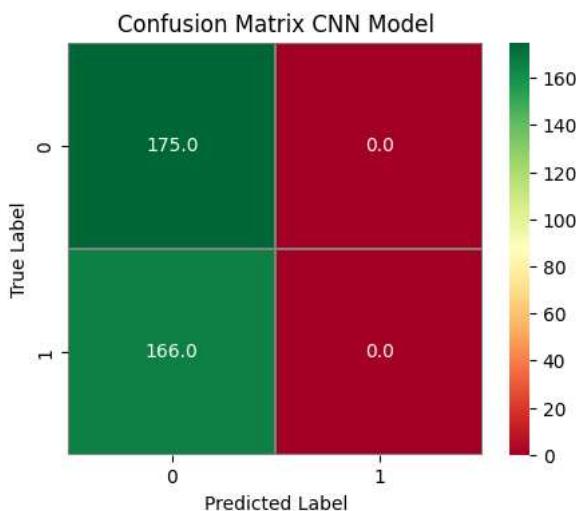
11/11 [=====] - 0s 22ms/step
Classification Report:
precision    recall   f1-score   support
          0       0.51      1.00      0.68     175
          1       0.00      0.00      0.00     166

   accuracy                           0.51    341
  macro avg       0.26      0.50      0.34    341
weighted avg       0.26      0.51      0.35    341

Confusion Matrix:
[[175  0]
 [166  0]]
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels without unique samples.
  _warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels without unique samples.
  _warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels without unique samples.
  _warn_prf(average, modifier, msg_start, len(result))

```

f,ax = plt.subplots(figsize=(5, 4))
sns.heatmap(cm0, annot=True, linewidths=0.01,cmap="RdYlGn",linecolor="gray", fmt= '.1f',ax=ax)
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.title("Confusion Matrix CNN Model")
plt.show()



CNN Second Iteration

```

model1 = models.Sequential()

model1.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(224, 224, 3)))
model1.add(layers.MaxPooling2D((2, 2)))
model1.add(layers.Conv2D(64, (3, 3), activation='relu'))
model1.add(layers.MaxPooling2D((2, 2)))
model1.add(layers.Conv2D(128, (3, 3), activation='relu'))
model1.add(layers.MaxPooling2D((2, 2)))

model1.add(layers.Flatten())
model1.add(layers.Dense(48, activation='relu'))
model1.add(layers.Dropout(0.5))
model1.add(layers.Dense(64, activation='relu')) # Additional layer
model1.add(layers.Dense(128, activation='relu')) # Additional layer
model1.add(layers.Dense(1, activation='sigmoid'))

model1.summary()

```

Model: "sequential"

```
=====
conv2d (Conv2D)           (None, 222, 222, 32)      896
max_pooling2d (MaxPooling2D) (None, 111, 111, 32)    0
conv2d_1 (Conv2D)          (None, 109, 109, 64)     18496
max_pooling2d_1 (MaxPooling2D) (None, 54, 54, 64)    0
conv2d_2 (Conv2D)          (None, 52, 52, 128)     73856
max_pooling2d_2 (MaxPooling2D) (None, 26, 26, 128)    0
flatten (Flatten)          (None, 86528)            0
dense (Dense)              (None, 48)                4153392
dropout (Dropout)          (None, 48)                0
dense_1 (Dense)             (None, 64)                3136
dense_2 (Dense)             (None, 128)               8320
dense_3 (Dense)             (None, 1)                 129
=====
```

```
Total params: 4258225 (16.24 MB)
Trainable params: 4258225 (16.24 MB)
Non-trainable params: 0 (0.00 Byte)
```

```
model1.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])

history1 = model1.fit(X_train, y_train, epochs=100, batch_size=32, validation_data=(X_val, y_val))

test_loss, test_acc = model1.evaluate(X_test, y_test)
print(f'Test accuracy: {test_acc}')
```

```
Epoch 94/100
38/38 [=====] - 2s 58ms/step - loss: 0.0234 - accuracy: 0.9857 - val_loss: 4.5368 - val_accuracy: 0.5893
Epoch 95/100
38/38 [=====] - 2s 59ms/step - loss: 0.0436 - accuracy: 0.9798 - val_loss: 3.9849 - val_accuracy: 0.6071
Epoch 96/100
38/38 [=====] - 2s 52ms/step - loss: 0.0424 - accuracy: 0.9832 - val_loss: 4.6946 - val_accuracy: 0.5952
Epoch 97/100
38/38 [=====] - 2s 51ms/step - loss: 0.0373 - accuracy: 0.9832 - val_loss: 3.9983 - val_accuracy: 0.6369
Epoch 98/100
38/38 [=====] - 2s 53ms/step - loss: 0.0330 - accuracy: 0.9874 - val_loss: 4.4225 - val_accuracy: 0.5952
Epoch 99/100
38/38 [=====] - 2s 52ms/step - loss: 0.0325 - accuracy: 0.9806 - val_loss: 3.9140 - val_accuracy: 0.6012
Epoch 100/100
38/38 [=====] - 2s 52ms/step - loss: 0.0225 - accuracy: 0.9840 - val_loss: 4.8677 - val_accuracy: 0.6131
11/11 [=====] - 1s 64ms/step - loss: 4.6356 - accuracy: 0.6276
Test accuracy: 0.6275659799575806
```

▼ Resnet50

```
from tensorflow.keras.applications import ResNet50
from tensorflow.keras import layers, models
from tensorflow.keras.optimizers import Adam

resnet50_model = ResNet50(weights='imagenet', include_top=False, input_shape=(224, 224, 3))

for layer in resnet50_model.layers:
    layer.trainable = False

model2 = models.Sequential()
model2.add(resnet50_model)
model2.add(layers.GlobalAveragePooling2D())
model2.add(layers.Dense(256, activation='relu'))
model2.add(layers.Dense(128, activation='relu'))
model2.add(layers.Dense(64, activation='relu'))
model2.add(layers.Dropout(0.5))
model2.add(layers.Dense(2, activation='sigmoid'))

model2.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
<hr/>		
resnet50 (Functional)	(None, 7, 7, 2048)	23587712
<hr/>		
global_average_pooling2d (GlobalAveragePooling2D)	(None, 2048)	0
dense (Dense)	(None, 256)	524544
dense_1 (Dense)	(None, 128)	32896
dense_2 (Dense)	(None, 64)	8256
dropout (Dropout)	(None, 64)	0
dense_3 (Dense)	(None, 2)	130
<hr/>		
Total params: 24153538 (92.14 MB)		
Trainable params: 565826 (2.16 MB)		
Non-trainable params: 23587712 (89.98 MB)		

```
model2.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])

history2 = model2.fit(X_train, y_train, epochs=50, batch_size=32, validation_data=(X_val, y_val))

test_loss, test_acc = model2.evaluate(X_test, y_test)
print(f'Test accuracy: {test_acc}')
```

```
Epoch 1/50
38/38 [=====] - 23s 200ms/step - loss: 0.7146 - accuracy: 0.4928 - val_loss: 0.6948 - val_accuracy: 0.4821
Epoch 2/50
38/38 [=====] - 4s 112ms/step - loss: 0.6953 - accuracy: 0.5274 - val_loss: 0.6930 - val_accuracy: 0.5179
Epoch 3/50
38/38 [=====] - 4s 108ms/step - loss: 0.6914 - accuracy: 0.5535 - val_loss: 0.6928 - val_accuracy: 0.5179
```

```

Epoch 4/50
38/38 [=====] - 4s 108ms/step - loss: 0.6922 - accuracy: 0.5375 - val_loss: 0.6925 - val_accuracy: 0.5179
Epoch 5/50
38/38 [=====] - 4s 114ms/step - loss: 0.6906 - accuracy: 0.5366 - val_loss: 0.6933 - val_accuracy: 0.5179
Epoch 6/50
38/38 [=====] - 4s 114ms/step - loss: 0.6910 - accuracy: 0.5484 - val_loss: 0.6957 - val_accuracy: 0.5179
Epoch 7/50
38/38 [=====] - 4s 109ms/step - loss: 0.6931 - accuracy: 0.5299 - val_loss: 0.6925 - val_accuracy: 0.5179
Epoch 8/50
38/38 [=====] - 4s 113ms/step - loss: 0.6934 - accuracy: 0.5333 - val_loss: 0.6934 - val_accuracy: 0.5179
Epoch 9/50
38/38 [=====] - 4s 116ms/step - loss: 0.6917 - accuracy: 0.5392 - val_loss: 0.6937 - val_accuracy: 0.5179
Epoch 10/50
38/38 [=====] - 4s 113ms/step - loss: 0.6906 - accuracy: 0.5484 - val_loss: 0.6949 - val_accuracy: 0.5179
Epoch 11/50
38/38 [=====] - 4s 109ms/step - loss: 0.6917 - accuracy: 0.5476 - val_loss: 0.6929 - val_accuracy: 0.5179
Epoch 12/50
38/38 [=====] - 4s 116ms/step - loss: 0.6902 - accuracy: 0.5476 - val_loss: 0.6927 - val_accuracy: 0.5179
Epoch 13/50
38/38 [=====] - 4s 118ms/step - loss: 0.6894 - accuracy: 0.5476 - val_loss: 0.6930 - val_accuracy: 0.5179
Epoch 14/50
38/38 [=====] - 4s 110ms/step - loss: 0.6899 - accuracy: 0.5476 - val_loss: 0.6933 - val_accuracy: 0.5179
Epoch 15/50
38/38 [=====] - 4s 110ms/step - loss: 0.6907 - accuracy: 0.5476 - val_loss: 0.6926 - val_accuracy: 0.5179
Epoch 16/50
38/38 [=====] - 4s 115ms/step - loss: 0.6892 - accuracy: 0.5476 - val_loss: 0.6930 - val_accuracy: 0.5179
Epoch 17/50
38/38 [=====] - 4s 116ms/step - loss: 0.6886 - accuracy: 0.5476 - val_loss: 0.6935 - val_accuracy: 0.5179
Epoch 18/50
38/38 [=====] - 4s 111ms/step - loss: 0.6877 - accuracy: 0.5476 - val_loss: 0.6943 - val_accuracy: 0.5179
Epoch 19/50
38/38 [=====] - 4s 114ms/step - loss: 0.6923 - accuracy: 0.5476 - val_loss: 0.6953 - val_accuracy: 0.5179
Epoch 20/50
38/38 [=====] - 5s 120ms/step - loss: 0.6912 - accuracy: 0.5476 - val_loss: 0.7002 - val_accuracy: 0.5179
Epoch 21/50
38/38 [=====] - 4s 111ms/step - loss: 0.6883 - accuracy: 0.5476 - val_loss: 0.6926 - val_accuracy: 0.5179
Epoch 22/50
38/38 [=====] - 4s 114ms/step - loss: 0.6892 - accuracy: 0.5476 - val_loss: 0.6937 - val_accuracy: 0.5179
Epoch 23/50
38/38 [=====] - 4s 117ms/step - loss: 0.6906 - accuracy: 0.5476 - val_loss: 0.6931 - val_accuracy: 0.5179
Epoch 24/50
38/38 [=====] - 4s 115ms/step - loss: 0.6892 - accuracy: 0.5476 - val_loss: 0.6940 - val_accuracy: 0.5179
Epoch 25/50
38/38 [=====] - 4s 114ms/step - loss: 0.6898 - accuracy: 0.5476 - val_loss: 0.6938 - val_accuracy: 0.5179
Epoch 26/50
38/38 [=====] - 4s 116ms/step - loss: 0.6890 - accuracy: 0.5476 - val_loss: 0.6934 - val_accuracy: 0.5179
Epoch 27/50
38/38 [=====] - 5s 121ms/step - loss: 0.6899 - accuracy: 0.5476 - val_loss: 0.6934 - val_accuracy: 0.5179
Epoch 28/50
38/38 [=====] - 4s 113ms/step - loss: 0.6891 - accuracy: 0.5476 - val_loss: 0.6939 - val_accuracy: 0.5179
Epoch 29/50
38/38 [=====] - 4s 112ms/step - loss: 0.6893 - accuracy: 0.5476 - val_loss: 0.6936 - val_accuracy: 0.5179

```

```

acc = history2.history['accuracy']
val_acc = history2.history['val_accuracy']
loss = history2.history['loss']
val_loss = history2.history['val_loss']

```

```
epochs = range(1, len(acc) + 1)
```

```

plt.plot(epochs, acc, 'r', label='Training acc')
plt.plot(epochs, val_acc, 'k', label='Validation acc')
plt.title('Resnet50 Training and Validation Accuracy')
plt.legend()

```

```
plt.figure()
```

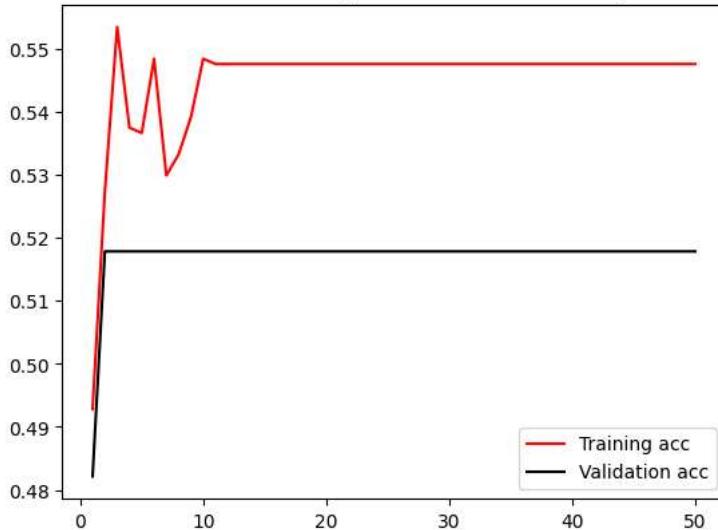
```

plt.plot(epochs, loss, 'r', label='Training loss')
plt.plot(epochs, val_loss, 'k', label='Validation loss')
plt.title('ResNet50 Training and Validation Loss')
plt.legend()

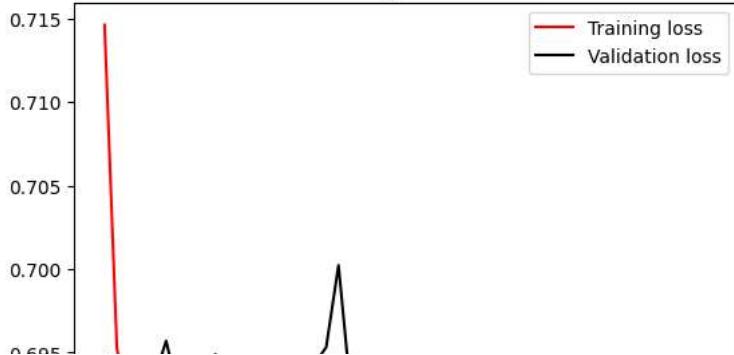
```

```
plt.show()
```

Resnet50 Training and Validation Accuracy



ResNet50 Training and Validation Loss



```
from sklearn.metrics import classification_report, confusion_matrix
```

```
y_pred = model2.predict(X_test)
y_pred_classes = np.argmax(y_pred, axis=1)
y_true_classes = np.argmax(y_test, axis=1)
```

```
report2 = classification_report(y_true_classes, y_pred_classes)
print("Classification Report:")
print(report2)
```

```
cm2 = confusion_matrix(y_true_classes, y_pred_classes)
print("Confusion Matrix:")
print(cm2)
```

```
11/11 [=====] - 3s 86ms/step
```

```
Classification Report:
```

	precision	recall	f1-score	support
0	0.51	1.00	0.68	175
1	0.00	0.00	0.00	166
accuracy			0.51	341
macro avg	0.26	0.50	0.34	341
weighted avg	0.26	0.51	0.35	341

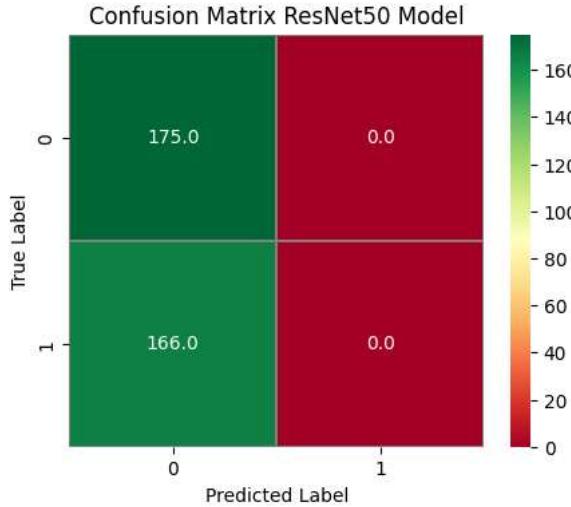
```
Confusion Matrix:
```

```
[[175  0]
 [166  0]]
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-defined and hence ignored! The labels must be binary or multiclass.
  _warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-defined and hence ignored! The labels must be binary or multiclass.
  _warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-defined and hence ignored! The labels must be binary or multiclass.
  _warn_prf(average, modifier, msg_start, len(result))
```

```
f,ax = plt.subplots(figsize=(5, 4))
sns.heatmap(cm2, annot=True, linewidths=0.01,cmap="RdYlGn",linecolor="gray", fmt= '.1f',ax=ax)
plt.xlabel("Predicted Label")
```

```
plt.ylabel("True Label")
plt.title("Confusion Matrix ResNet50 Model")
plt.show()
```



ResNet50 Second Iteration

```
resnet50_model = ResNet50(weights='imagenet', include_top=False, input_shape=(224, 224, 3))
```

```
for layer in resnet50_model.layers:
    layer.trainable = False
```

```
model3 = models.Sequential()
model3.add(resnet50_model)
model3.add(layers.GlobalAveragePooling2D())
model3.add(layers.Dense(256, activation='relu'))
model3.add(layers.Dropout(0.5))
model3.add(layers.Dense(128, activation='relu'))
model3.add(layers.Dense(64, activation='relu'))
model3.add(layers.Dense(32, activation='relu'))
model3.add(layers.Dense(16, activation='relu'))
model3.add(layers.Dense(1, activation='softmax'))
```

```
model3.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
<hr/>		
resnet50 (Functional)	(None, 7, 7, 2048)	23587712
<hr/>		
global_average_pooling2d (GlobalAveragePooling2D)	(None, 2048)	0
dense (Dense)	(None, 256)	524544
dropout (Dropout)	(None, 256)	0
dense_1 (Dense)	(None, 128)	32896
dense_2 (Dense)	(None, 64)	8256
dense_3 (Dense)	(None, 32)	2080
dense_4 (Dense)	(None, 16)	528
dense_5 (Dense)	(None, 2)	34
<hr/>		
Total params: 24156050 (92.15 MB)		
Trainable params: 568338 (2.17 MB)		
Non-trainable params: 23587712 (89.98 MB)		

```
model3.compile(optimizer=Adam(learning_rate=0.0001), loss='categorical_crossentropy', metrics=['accuracy'])
```

```

history3 = model3.fit(X_train, y_train, epochs=100, batch_size=32, validation_data=(X_val, y_val))

test_loss, test_acc = model3.evaluate(X_test, y_test)
print(f'Test accuracy: {test_acc}')

Epoch 47/100
38/38 [=====] - 5s 125ms/step - loss: 0.6885 - accuracy: 0.5459 - val_loss: 0.6928 - val_accuracy: 0.5179
Epoch 48/100
38/38 [=====] - 8s 223ms/step - loss: 0.6889 - accuracy: 0.5451 - val_loss: 0.6930 - val_accuracy: 0.5179
Epoch 49/100
38/38 [=====] - 5s 126ms/step - loss: 0.6901 - accuracy: 0.5484 - val_loss: 0.6929 - val_accuracy: 0.5179
Epoch 50/100
38/38 [=====] - 4s 119ms/step - loss: 0.6897 - accuracy: 0.5484 - val_loss: 0.6930 - val_accuracy: 0.5179
Epoch 51/100
38/38 [=====] - 4s 116ms/step - loss: 0.6898 - accuracy: 0.5476 - val_loss: 0.6928 - val_accuracy: 0.5179
Epoch 52/100
38/38 [=====] - 4s 117ms/step - loss: 0.6894 - accuracy: 0.5476 - val_loss: 0.6931 - val_accuracy: 0.5179
Epoch 53/100
38/38 [=====] - 5s 122ms/step - loss: 0.6893 - accuracy: 0.5468 - val_loss: 0.6925 - val_accuracy: 0.5179
Epoch 54/100
38/38 [=====] - 4s 119ms/step - loss: 0.6891 - accuracy: 0.5476 - val_loss: 0.6926 - val_accuracy: 0.5179
Epoch 55/100
38/38 [=====] - 4s 118ms/step - loss: 0.6903 - accuracy: 0.5459 - val_loss: 0.6923 - val_accuracy: 0.5179
Epoch 56/100
38/38 [=====] - 8s 224ms/step - loss: 0.6898 - accuracy: 0.5451 - val_loss: 0.6929 - val_accuracy: 0.5179
Epoch 57/100
38/38 [=====] - 4s 114ms/step - loss: 0.6902 - accuracy: 0.5451 - val_loss: 0.6928 - val_accuracy: 0.5179
Epoch 58/100
38/38 [=====] - 9s 229ms/step - loss: 0.6889 - accuracy: 0.5476 - val_loss: 0.6927 - val_accuracy: 0.5179
Epoch 59/100
38/38 [=====] - 4s 117ms/step - loss: 0.6896 - accuracy: 0.5468 - val_loss: 0.6926 - val_accuracy: 0.5179
Epoch 60/100
38/38 [=====] - 6s 162ms/step - loss: 0.6882 - accuracy: 0.5468 - val_loss: 0.6930 - val_accuracy: 0.5179
Epoch 61/100
38/38 [=====] - 5s 121ms/step - loss: 0.6885 - accuracy: 0.5476 - val_loss: 0.6932 - val_accuracy: 0.5179
Epoch 62/100
38/38 [=====] - 4s 114ms/step - loss: 0.6873 - accuracy: 0.5510 - val_loss: 0.6931 - val_accuracy: 0.5179
Epoch 63/100
38/38 [=====] - 4s 115ms/step - loss: 0.6895 - accuracy: 0.5484 - val_loss: 0.6936 - val_accuracy: 0.5179
Epoch 64/100
38/38 [=====] - 5s 122ms/step - loss: 0.6908 - accuracy: 0.5476 - val_loss: 0.6928 - val_accuracy: 0.5179
Epoch 65/100
38/38 [=====] - 5s 121ms/step - loss: 0.6897 - accuracy: 0.5476 - val_loss: 0.6924 - val_accuracy: 0.5179
Epoch 66/100
38/38 [=====] - 4s 116ms/step - loss: 0.6886 - accuracy: 0.5476 - val_loss: 0.6925 - val_accuracy: 0.5179
Epoch 67/100
38/38 [=====] - 5s 119ms/step - loss: 0.6886 - accuracy: 0.5484 - val_loss: 0.6925 - val_accuracy: 0.5179
Epoch 68/100
38/38 [=====] - 5s 122ms/step - loss: 0.6870 - accuracy: 0.5493 - val_loss: 0.6929 - val_accuracy: 0.5179
Epoch 69/100
38/38 [=====] - 4s 119ms/step - loss: 0.6893 - accuracy: 0.5484 - val_loss: 0.6932 - val_accuracy: 0.5179
Epoch 70/100
38/38 [=====] - 4s 116ms/step - loss: 0.6879 - accuracy: 0.5484 - val_loss: 0.6931 - val_accuracy: 0.5179
Epoch 71/100
38/38 [=====] - 184s 5s/step - loss: 0.6893 - accuracy: 0.5501 - val_loss: 0.6931 - val_accuracy: 0.5179
Epoch 72/100
38/38 [=====] - 87s 2s/step - loss: 0.6908 - accuracy: 0.5468 - val_loss: 0.6930 - val_accuracy: 0.5179
Epoch 73/100
38/38 [=====] - 7s 184ms/step - loss: 0.6888 - accuracy: 0.5468 - val_loss: 0.6924 - val_accuracy: 0.5179
Epoch 74/100
38/38 [=====] - 4s 116ms/step - loss: 0.6890 - accuracy: 0.5484 - val_loss: 0.6932 - val_accuracy: 0.5179
Epoch 75/100
38/38 [=====] - 4s 119ms/step - loss: 0.6892 - accuracy: 0.5484 - val_loss: 0.6931 - val_accuracy: 0.5179
Epoch 76/100

```

▼ CNN Model VGG16

```

from keras.applications import vgg16

vgg16_model = tf.keras.applications.vgg16.VGG16(
    weights = 'imagenet',
    include_top = False,
    input_shape = (224, 224, 3)
)

for layer in vgg16_model.layers:
    layer.trainable = False

model4 = tf.keras.Sequential()

```

```
model4.add(vgg16_model)
model4.add(tf.keras.layers.GlobalAveragePooling2D())
model4.add(tf.keras.layers.Dropout(0.5))
model4.add(tf.keras.layers.Dense(256, activation = 'relu'))
model4.add(tf.keras.layers.Dense(128, activation = 'relu'))
model4.add(tf.keras.layers.Dense(64, activation = 'relu'))
model4.add(tf.keras.layers.Dense(2, activation = 'softmax'))
```

```
model4.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
<hr/>		
vgg16 (Functional)	(None, 7, 7, 512)	14714688
<hr/>		
global_average_pooling2d_1 (GlobalAveragePooling2D)	(None, 512)	0
dropout_1 (Dropout)	(None, 512)	0
dense_4 (Dense)	(None, 256)	131328
dense_5 (Dense)	(None, 128)	32896
dense_6 (Dense)	(None, 64)	8256
dense_7 (Dense)	(None, 2)	130
<hr/>		
Total params: 14887298 (56.79 MB)		
Trainable params: 172610 (674.26 KB)		
Non-trainable params: 14714688 (56.13 MB)		

```
model4.compile(optimizer = 'adam', loss='binary_crossentropy',metrics=['accuracy'])

history4 = model4.fit(X_train, y_train, epochs = 50, batch_size = 32, validation_data=(X_val, y_val))

test_loss, test_acc = model4.evaluate(X_test, y_test)
print(f'Test accuracy: {test_acc}')
```

```
38/38 [=====] - 6s 16ms/step - loss: 0.6446 - accuracy: 0.6552 - val_loss: 0.6272 - val_accuracy: 0.6190
Epoch 43/50
38/38 [=====] - 6s 169ms/step - loss: 0.6418 - accuracy: 0.6276 - val_loss: 0.6185 - val_accuracy: 0.6369
Epoch 44/50
38/38 [=====] - 6s 151ms/step - loss: 0.6459 - accuracy: 0.6259 - val_loss: 0.6288 - val_accuracy: 0.6012
Epoch 45/50
38/38 [=====] - 6s 152ms/step - loss: 0.6448 - accuracy: 0.6369 - val_loss: 0.6307 - val_accuracy: 0.6548
Epoch 46/50
38/38 [=====] - 7s 177ms/step - loss: 0.6429 - accuracy: 0.6327 - val_loss: 0.6238 - val_accuracy: 0.6488
Epoch 47/50
38/38 [=====] - 7s 174ms/step - loss: 0.6369 - accuracy: 0.6369 - val_loss: 0.6271 - val_accuracy: 0.6250
Epoch 48/50
38/38 [=====] - 7s 181ms/step - loss: 0.6440 - accuracy: 0.6293 - val_loss: 0.6252 - val_accuracy: 0.6488
Epoch 49/50
38/38 [=====] - 8s 199ms/step - loss: 0.6449 - accuracy: 0.6209 - val_loss: 0.6542 - val_accuracy: 0.6429
Epoch 50/50
38/38 [=====] - 6s 151ms/step - loss: 0.6432 - accuracy: 0.6175 - val_loss: 0.6254 - val_accuracy: 0.6548
11/11 [=====] - 4s 379ms/step - loss: 0.6306 - accuracy: 0.6598
Test accuracy: 0.6598240733146667
```

```
acc = history4.history['accuracy']
val_acc = history4.history['val_accuracy']
loss = history4.history['loss']
val_loss = history4.history['val_loss']

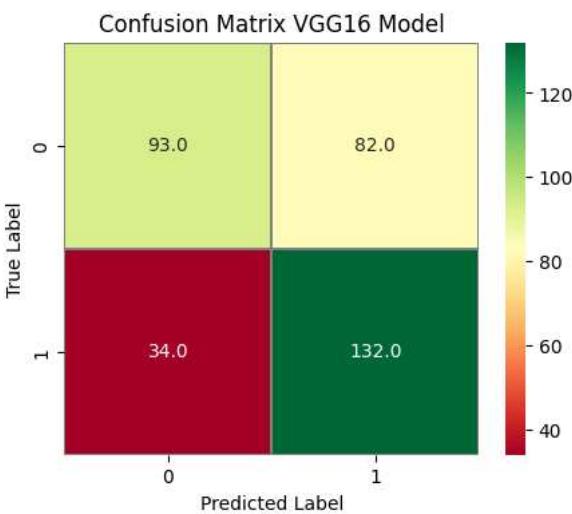
epochs = range(1, len(acc) + 1)

plt.plot(epochs, acc, 'r', label='Training acc')
plt.plot(epochs, val_acc, 'g', label='Validation acc')
plt.title('VGG16 Training and Validation Accuracy')
plt.legend()

plt.figure()

plt.plot(epochs, loss, 'r', label='Training loss')
plt.plot(epochs, val_loss, 'g', label='Validation loss')
plt.title('VGG16 Training and Validation Loss')
plt.legend()

plt.show()
```



Second Iteration of VGG16 Model

```
vgg16_model = tf.keras.applications.vgg16.VGG16()
    weights = 'imagenet',
    include_top = False,
    input_shape = (224, 224, 3)
)
for layer in vgg16_model.layers:
```

```
layer.trainable = False

model5 = tf.keras.Sequential()
model5.add(vgg16_model)
model5.add(tf.keras.layers.GlobalAveragePooling2D())
model5.add(tf.keras.layers.Dropout(0.5))
model5.add(tf.keras.layers.Dense(256, activation = 'relu'))
model5.add(tf.keras.layers.Dense(128, activation='relu'))
model5.add(tf.keras.layers.Dense(64, activation='relu'))
model5.add(tf.keras.layers.Dense(32, activation='relu'))
model5.add(tf.keras.layers.Dense(16, activation='relu'))
model5.add(tf.keras.layers.Dense(2, activation = 'sigmoid'))

model5.summary()

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16\_weights\_tf\_dim\_ordering\_tf\_kernels\_notop.h5
58889256/58889256 [=====] - 0s 0us/step
Model: "sequential"

Layer (type)          Output Shape         Param #
=================================================================
vgg16 (Functional)    (None, 7, 7, 512)     14714688
global_average_pooling2d ( GlobalAveragePooling2D)
                                         (None, 512)           0
dropout (Dropout)      (None, 512)           0
dense (Dense)          (None, 256)          131328
dense_1 (Dense)        (None, 128)          32896
dense_2 (Dense)        (None, 64)           8256
dense_3 (Dense)        (None, 32)           2080
dense_4 (Dense)        (None, 16)           528
dense_5 (Dense)        (None, 2)            34
=====
Total params: 14889810 (56.80 MB)
Trainable params: 175122 (684.07 KB)
Non-trainable params: 14714688 (56.13 MB)
```

```
model5.compile(optimizer = 'rmsprop', loss='binary_crossentropy',metrics=['accuracy'])

history5 = model5.fit(X_train, y_train, epochs = 50, batch_size = 32, validation_data=(X_val, y_val))

test_loss, test_acc = model5.evaluate(X_test, y_test)
print(f'Test accuracy: {test_acc}')
```

```

38/38 [=====] - 6s 170ms/step - loss: 0.6684 - accuracy: 0.6091 - val_loss: 0.6552 - val_accuracy: 0.6429
Epoch 37/50
38/38 [=====] - 6s 167ms/step - loss: 0.6711 - accuracy: 0.5729 - val_loss: 0.6530 - val_accuracy: 0.6012
Epoch 38/50
38/38 [=====] - 6s 168ms/step - loss: 0.6688 - accuracy: 0.5973 - val_loss: 0.6483 - val_accuracy: 0.6369
Epoch 39/50
38/38 [=====] - 6s 170ms/step - loss: 0.6652 - accuracy: 0.5981 - val_loss: 0.6581 - val_accuracy: 0.5893
Epoch 40/50
38/38 [=====] - 6s 150ms/step - loss: 0.6738 - accuracy: 0.5796 - val_loss: 0.6695 - val_accuracy: 0.6250
Epoch 41/50
38/38 [=====] - 6s 155ms/step - loss: 0.6645 - accuracy: 0.6040 - val_loss: 0.6551 - val_accuracy: 0.6250
Epoch 42/50
38/38 [=====] - 6s 153ms/step - loss: 0.6570 - accuracy: 0.6259 - val_loss: 0.6432 - val_accuracy: 0.6071
Epoch 43/50
38/38 [=====] - 6s 168ms/step - loss: 0.6700 - accuracy: 0.5906 - val_loss: 0.6596 - val_accuracy: 0.5476
Epoch 44/50
38/38 [=====] - 6s 156ms/step - loss: 0.6578 - accuracy: 0.6158 - val_loss: 0.6539 - val_accuracy: 0.6250
Epoch 45/50
38/38 [=====] - 6s 169ms/step - loss: 0.6634 - accuracy: 0.6226 - val_loss: 0.6463 - val_accuracy: 0.6131
Epoch 46/50
38/38 [=====] - 6s 167ms/step - loss: 0.6671 - accuracy: 0.5838 - val_loss: 0.6642 - val_accuracy: 0.5298
Epoch 47/50
38/38 [=====] - 6s 171ms/step - loss: 0.6705 - accuracy: 0.6024 - val_loss: 0.6595 - val_accuracy: 0.6310
Epoch 48/50
38/38 [=====] - 6s 167ms/step - loss: 0.6635 - accuracy: 0.6108 - val_loss: 0.6477 - val_accuracy: 0.6071
Epoch 49/50
38/38 [=====] - 6s 168ms/step - loss: 0.6599 - accuracy: 0.6099 - val_loss: 0.6571 - val_accuracy: 0.6429
Epoch 50/50
38/38 [=====] - 6s 156ms/step - loss: 0.6606 - accuracy: 0.6108 - val_loss: 0.6370 - val_accuracy: 0.6429
11/11 [=====] - 4s 388ms/step - loss: 0.6385 - accuracy: 0.6657
Test accuracy: 0.6656891703605652

```

```

acc = history5.history['accuracy']
val_acc = history5.history['val_accuracy']
loss = history5.history['loss']
val_loss = history5.history['val_loss']

epochs = range(1, len(acc) + 1)

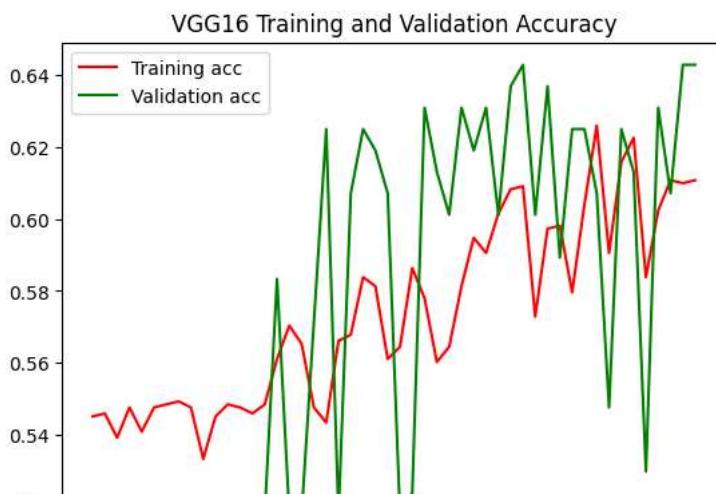
plt.plot(epochs, acc, 'r', label='Training acc')
plt.plot(epochs, val_acc, 'g', label='Validation acc')
plt.title('VGG16 Training and Validation Accuracy')
plt.legend()

plt.figure()

plt.plot(epochs, loss, 'r', label='Training loss')
plt.plot(epochs, val_loss, 'g', label='Validation loss')
plt.title('VGG16 Training and Validation Loss')
plt.legend()

plt.show()

```



```
from sklearn.metrics import classification_report, confusion_matrix
```

```
y_pred = model5.predict(X_test)
y_pred_classes = np.argmax(y_pred, axis=1)
y_true_classes = np.argmax(y_test, axis=1)

report5 = classification_report(y_true_classes, y_pred_classes)
print("Classification Report:")
print(report5)

cm5 = confusion_matrix(y_true_classes, y_pred_classes)
print("Confusion Matrix:")
print(cm5)
```

```
11/11 [=====] - 2s 116ms/step
Classification Report:
precision    recall    f1-score   support
          0       0.67      0.69      0.68     175
          1       0.66      0.64      0.65     166
   accuracy                           0.67     341
  macro avg       0.67      0.66      0.67     341
weighted avg       0.67      0.67      0.67     341
```

```
Confusion Matrix:
[[121  54]
 [ 60 106]]
```

```
f,ax = plt.subplots(figsize=(5, 4))
sns.heatmap(cm5, annot=True, linewidths=0.01,cmap="RdYlGn",linecolor="gray", fmt= '.1f',ax=ax)
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.title("Confusion Matrix VGG16 Model")
plt.show()
```

